



**Spark
Playground**

PySpark Syntax Cheat Sheet - Quick Syntax Guide For Data Engineers

Welcome to the **Spark Playground's PySpark Syntax Cheat Sheet**! This document is designed to provide a quick reference to essential PySpark functions and operations.

Check out the Spark Playground website [here](#)

Create DataFrame

With Default Schema

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Example").getOrCreate()

data = [(1, "Alice", 29), (2, "Bob", 35)]
df = spark.createDataFrame(data, ["id", "name", "age"])
df.show()
```

Explicit Schema

```
from pyspark.sql.types import StructType, StructField, IntegerType, StringType
```

```
schema = StructType([  
    StructField("id", IntegerType(), True),  
    StructField("name", StringType(), True),  
    StructField("age", IntegerType(), True)  
])
```

```
df = spark.createDataFrame(data, schema)  
df.printSchema()  
df.show()
```

```
# Schema as a string
```

```
data = [(1, "Alice", 29), (2, "Bob", 35)]  
schema = "id INT, name STRING, age INT"  
df = spark.createDataFrame(data, schema=schema)
```

```
# Schema String with Float and Boolean Types
```

```
schema = "id INT, name STRING, salary FLOAT, is_active BOOLEAN"  
data = [(1, "Alice", 50000.75, True), (2, "Bob", 60000.50, False)]  
df = spark.createDataFrame(data, schema=schema)
```

```
# Schema String with Date and Timestamp
```

```
from datetime import date, datetime  
schema = "id INT, name STRING, join_date DATE, last_login TIMESTAMP"  
data = [(1, "Alice", date(2023, 1, 15), datetime(2024, 3, 10, 14, 30, 0)),  
        (2, "Bob", date(2023, 1, 15), datetime(2024, 3, 10, 14, 30, 0))]  
df = spark.createDataFrame(data, schema=schema)
```

Using a List of Dictionaries

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Example").getOrCreate()

data = [
    {"id": 1, "name": "Alice", "age": 29},
    {"id": 2, "name": "Bob", "age": 35}
]

df = spark.createDataFrame(data)
df.show()
```

Reading Files

CSV Files

```
#Basic CSV files
df = spark.read.format("csv").load("/path/to/sample.csv")

#csv with header
df = spark.read.option("header", True).csv("/path/to/sample.csv")

# multiple options
df = spark.read.option("inferSchema", True).option("delimiter", "," ).csv("/path/to/sample.csv")

# with defined schema
from pyspark.sql.types import StructType, StructField, StringType, IntegerType
schema = StructType([
    StructField("name", StringType(), True),
    StructField("age", IntegerType(), True)
])
df = spark.read.format("csv").schema(schema).load("/path/to/sample.csv")
```

JSON Files

```
# Basic JSON file
df = spark.read.format("json").load("/path/to/sample.json")

# JSON with multi-line records
df = spark.read.option("multiline", True).json("/path/to/sample.json")

# JSON with a defined schema
from pyspark.sql.types import StructType, StructField, StringType, IntegerType

schema = StructType([
    StructField("name", StringType(), True),
    StructField("age", IntegerType(), True)
])

df = spark.read.format("json").schema(schema).load("/path/to/sample.json")
```

Here's the **full set of examples for writing files in PySpark**, covering **CSV, JSON, Parquet, ORC, and Delta** formats:

Writing Files

CSV Files

Basic write to CSV

```
df.write.csv("/path/to/output_csv")
```

With header

```
df.write.option("header", True).csv("/path/to/output_csv")
```

With multiple options

```
df.write.option("header", True)\  
  .option("delimiter", ",")\  
  .option("quote", "\"")\  
  .csv("/path/to/output_csv")
```

Overwrite existing files

```
df.write.mode("overwrite").option("header", True).csv("/path/to/output_csv")
```

Append to existing data

```
df.write.mode("append").option("header", True).csv("/path/to/output_csv")
```

Write as a single file

```
df.coalesce(1).write.option("header", True).csv("/path/to/output_csv")
```

JSON Files

Basic JSON write

```
df.write.json("/path/to/output_json")
```

Overwrite mode

```
df.write.mode("overwrite").json("/path/to/output_json")
```

Append mode

```
df.write.mode("append").json("/path/to/output_json")
```

Pretty format (for readability)

```
df.write.option("compression", "none").json("/path/to/output_json")
```

Partitioned output

```
df.write.partitionBy("column_name").json("/path/to/output_json")
```

Parquet Files

```
# Basic Parquet write
df.write.parquet("/path/to/output_parquet")

# Overwrite mode
df.write.mode("overwrite").parquet("/path/to/output_parquet")

# Append mode
df.write.mode("append").parquet("/path/to/output_parquet")

# Partitioned output
df.write.partitionBy("column_name").parquet("/path/to/output_parquet")

# Compression options (default is snappy)
df.write.option("compression", "gzip").parquet("/path/to/output_parquet")
```

ORC Files

```
# Basic ORC write
df.write.orc("/path/to/output_orc")

# Overwrite mode
df.write.mode("overwrite").orc("/path/to/output_orc")

# Append mode
df.write.mode("append").orc("/path/to/output_orc")

# Partitioned output
df.write.partitionBy("column_name").orc("/path/to/output_orc")

# Compression options
df.write.option("compression", "zlib").orc("/path/to/output_orc")
```

Delta Tables/Files (Requires Delta Lake)

```
# Basic Delta write
df.write.format("delta").save("/path/to/output_delta")

# Basic Delta Table
df.write.format("delta").saveAsTable("table_name")

# Overwrite mode
df.write.format("delta").mode("overwrite").save("/path/to/output_delta")

# Append mode
df.write.format("delta").mode("append").save("/path/to/output_delta")

# With Delta-specific options (mergeSchema, overwriteSchema)
df.write.format("delta")\
  .option("overwriteSchema", "true")\
  .mode("overwrite")\
  .save("/path/to/output_delta")

# Partitioned output
df.write.format("delta").partitionBy("column_name").save("/path/to/output_delta")

# Bucketing
df.write.bucketBy(numBuckets=4, col="column_name")\
  .sortBy("column_name")\
  .saveAsTable("table_name")

# Bucketing by "name" into 4 buckets, sorting within each bucket
df.write \
  .bucketBy(4, "name") \
  .sortBy("age") \
  .mode("overwrite") \
  .saveAsTable("bucketed_people")
```

Select, Drop, Rename Columns

Selecting Columns


```
# Select single column
df = df.select("name")

# Select multiple columns
df = df.select("name", "age")

# Select columns dynamically
columns_to_select = ["name", "department"]
df = df.select(*columns_to_select)
```

Renaming Columns

```
# Rename a column
df = df.withColumnRenamed("name", "full_name")

# Rename multiple columns with chained calls
df = df.withColumnRenamed("old_col1", "new_col1")\
    .withColumnRenamed("old_col2", "new_col2")

# Rename columns using select and alias
from pyspark.sql.functions import col
df = df.select(
    col("old_column_name1").alias("new_column_name1"),
    col("old_column_name2").alias("new_column_name2"),
    # Add more columns as needed
)
```

Adding Columns

```
from pyspark.sql.functions import col, lit, expr, when
```

```
# Add a new column with a constant value
```

```
df = df.withColumn("country", lit("USA"))
```

```
# Add a new column with a calculated value
```

```
df = df.withColumn("salary_after_bonus", col("salary") * 1.1)
```

```
# Add a column using an SQL expression
```

```
df = df.withColumn("tax", expr("salary * 0.2"))
```

```
# Add a column with conditional logic
```

```
df = df.withColumn("high_earner", when(col("salary") > 55000, "Yes").otherwise("No"))
```

```
# Case When with multiple conditions
```

```
df = df.withColumn(  
    "salary_category",  
    when(col("salary") < 60000, "Low")  
    .when((col("salary") >= 60000) & (col("salary") < 90000), "Medium")  
    .otherwise("High")  
)
```

```
# Add multiple columns at once
```

```
df = df.withColumns({  
    "bonus": col("salary") * 0.1,  
    "net_salary": col("salary") - (col("salary") * 0.2)  
})
```

Dropping Columns

```
# Drop a column
```

```
df = df.drop("department")
```

```
# Drop multiple columns
```

```
df = df.drop('column1', 'column2', 'column3')
```

Filtering

You can refer to columns using any of these notations: `df.age` , `df['age']` , `col('age')`

Basic Filtering

```
# Filter on >, <, >=, <=, == condition
df_filtered = df.filter(df.age > 30)
df_filtered = df.filter(df['age'] > 30)

# Using col() function
from pyspark.sql.functions import col
df_filtered = df.filter(col("age") > 30)
```

Filter with Multiple Conditions

```
# Multiple conditions require parentheses around each condition

# AND condition ( & )
df_filtered = df.filter((df.age > 25) & (df.department == "Engineering"))
# OR condition ( | )
df_filtered = df.filter((df.age < 30) | (df.department == "Finance"))
```

String Filters

```
# Filter rows where department equals 'Marketing'
df_filtered = df.filter(df.department == "Marketing")

# Case-insensitive filter
df_filtered = df.filter(col("department").like("MARKETING"))

# Contains a substring
df_filtered = df.filter(col("department").contains("Engineer"))

# Filter rows where the name starts with 'A'
df.filter(col("name").startswith("A")).show()

# Filter rows where the name ends with 'e'
df.filter(col("name").endswith("e")).show()

# Filter rows where the name matches a regex
df.filter(col("name").rlike("^A.*")).show()
```

Null Filters

```
# Filter rows where a column is null
df_filtered = df.filter(df.department.isNull())
# Filter rows where a column is not null
df_filtered = df.filter(df.department.isNotNull())
```

Filter from a List

```
# Filter rows where department is in a list
departments = ["Engineering", "Finance"]
df_filtered = df.filter(col("department").isin(departments))
# Negate the filter (not in list)
df_filtered = df.filter(~col("department").isin(departments))
```

Data Cleansing

```
# 1. Drop all fully duplicate rows  
# Removes rows where all columns match exactly  
df = df.dropDuplicates()  
  
# 2. Drop duplicates based on specific columns  
# Keeps the first row for each unique email  
df = df.dropDuplicates(["email"])  
  
# 3. Get only distinct rows (same as SELECT DISTINCT)  
# Removes duplicates across all columns  
df = df.distinct()  
  
# 4. Drop rows with any null values  
# Removes rows with even a single null field  
df = df.dropna()  
  
# 5. Drop rows with nulls in specific columns  
# Only keeps rows where 'email' and 'age' are not null  
df = df.dropna(subset=["email", "age"])  
  
# 6. Fill missing values for all columns  
# Replaces all nulls with a default value  
df = df.fillna("N/A")  
  
# 7. Fill missing values for specific columns  
# Sets default age as 0 and country as "Unknown" if missing  
df = df.fillna({"age": 0, "country": "Unknown"})
```

Grouping

Import the required functions

```
from pyspark.sql.functions import count, sum, avg, min, max, countDistinct, col
```

Basic Aggregations without Grouping

```
#Count rows
df.count()

#Count Distinct Values in a column
df.select(countDistinct("Department")).show()

#Sum
df.select(sum("Salary")).show()

#Multiple Aggregations
df.select(min("Salary"), max("Salary")).show()
```

Aggregations with Grouping

```
#Group by a single column
df.groupBy("Department").sum("Salary").show()

#GroupBy with Multiple Columns
df.groupBy("Department", "Employee").sum("Salary").show()

#Group by with multiple aggregations
df.groupBy("Department").agg(
    count("Employee").alias("Employee_Count"),
    avg("Salary").alias("Average_Salary"),
    max("Salary").alias("Max_Salary")
)

#Filter after aggregation
df.groupBy("Department").agg(sum("Salary").alias("Total_Salary")).filter("Total_Sala
```

Common Aggregation Functions

Function	Description	Example
<code>count()</code>	Counts rows in a group.	<code>groupBy("Department").count()</code>
<code>sum()</code>	Sums values in a group.	<code>groupBy("Department").sum("Salary")</code>
<code>avg() / mean()</code>	Calculates average values.	<code>groupBy("Department").avg("Salary")</code>
<code>min()</code>	Finds the minimum value.	<code>groupBy("Department").min("Salary")</code>
<code>max()</code>	Finds the maximum value.	<code>groupBy("Department").max("Salary")</code>
<code>countDistinct()</code>	Counts distinct values in a group.	<code>countDistinct("Employee")</code>
<code>collect_list()</code>	Collects all values into a list.	<code>collect_list("Employee")</code>
<code>collect_set()</code>	Collects unique values into a set.	<code>collect_set("Employee")</code>

Joins

Join Types

Join Type	Syntax	Description
inner	<code>how="inner"</code>	Returns matching rows from both DataFrames based on the join condition.
outer (full)	<code>how="outer"</code>	Returns all rows, with <code>NULL</code> where no match is found in either DataFrame.
left (left_outer)	<code>how="left"</code>	Returns all rows from the left DataFrame, with <code>NULL</code> for unmatched rows in the right.
right (right_outer)	<code>how="right"</code>	Returns all rows from the right DataFrame, with <code>NULL</code> for unmatched rows in the left.
left_semi	<code>how="left_semi"</code>	This is just an inner join of the two DataFrames, but only returns columns of left DataFrame.
left_anti	<code>how="left_anti"</code>	Returns rows from the left DataFrame that do not have a match in the right.

Join Type	Syntax	Description
cross	<code>df1.crossJoin(df2)</code>	Returns the Cartesian product of rows from both DataFrames (no join condition).

Basic Syntax

```

# Basic Join
df1.join(df2, on="id", how="inner")

# Join on Multiple Columns
df1.join(df2, on=["col1", "col2"], how="left")

# Conditional Join
df1.join(df2, (df1.id == df2.id) & (df2.city == "New York"), how="inner")
# Multiple join conditions require parentheses around each condition
joined_df = sales_df.join(
    customers_df,
    (sales_df["customer_id"] == customers_df["customer_id"]) & (sales_df["region"] =
"inner"
)

# Select ALL columns from df1, and SOME columns from df2 (useful for left joins)
result = df1.join(df2, on="id", how="left").select(df1["*"], df2["state"], df2["town"])

# Broadcast Join for Small DataFrames
from pyspark.sql.functions import broadcast
df1.join(broadcast(df2), on="id", how="inner")

```

Date and Time Functions

String to Date Format


```
from pyspark.sql.functions import to_date
```

```
# 1. Convert string date to date type (using "yyyy-MM-dd")
```

```
# Input: "2025-01-25"
```

```
# Output: 2025-01-25 (as a Date type)
```

```
df = df.withColumn("date_parsed1", to_date("date_str", "yyyy-MM-dd"))
```

```
# 2. Convert string date to date type (using "dd-MMM-yyyy")
```

```
# Input: "25-Jan-2025"
```

```
# Output: 2025-01-25 (as a Date type)
```

```
df = df.withColumn("date_parsed2", to_date("date_str", "dd-MMM-yyyy"))
```

```
# 3. Convert string date to date type (using "MM/dd/yyyy")
```

```
# Input: "01/25/2025"
```

```
# Output: 2025-01-25 (as a Date type)
```

```
df = df.withColumn("date_parsed3", to_date("date_str", "MM/dd/yyyy"))
```

```
# 4. Convert string date to date type (using "yyyy.MM.dd")
```

```
# Input: "2025.01.25"
```

```
# Output: 2025-01-25 (as a Date type)
```

```
df = df.withColumn("date_parsed4", to_date("date_str", "yyyy.MM.dd"))
```

String to Timestamp Format

```
from pyspark.sql.functions import to_timestamp
```

```
# 1. Convert string timestamp to timestamp type (using "yyyy-MM-dd HH:mm:ss")
```

```
# Input: "2025-01-25 10:15:00"
```

```
# Output: 2025-01-25 10:15:00 (as a Timestamp type)
```

```
df = df.withColumn("timestamp_parsed1", to_timestamp("timestamp_str", "yyyy-MM-dd HH:mm:ss"))
```

```
# 2. Convert string timestamp to timestamp type (using "dd-MMM-yyyy HH:mm:ss")
```

```
# Input: "25-Jan-2025 10:15:00"
```

```
# Output: 2025-01-25 10:15:00 (as a Timestamp type)
```

```
df = df.withColumn("timestamp_parsed2", to_timestamp("timestamp_str", "dd-MMM-yyyy HH:mm:ss"))
```

```
# 3. Convert string timestamp to timestamp type (using "MM/dd/yyyy HH:mm:ss")
```

```
# Input: "01/25/2025 10:15:00"
```

```
# Output: 2025-01-25 10:15:00 (as a Timestamp type)
```

```
df = df.withColumn("timestamp_parsed3", to_timestamp("timestamp_str", "MM/dd/yyyy HH:mm:ss"))
```

```
# 4. Convert string timestamp to timestamp type (using "yyyy.MM.dd HH:mm:ss")
```

```
# Input: "2025.01.25 10:15:00"
```

```
# Output: 2025-01-25 10:15:00 (as a Timestamp type)
```

```
df = df.withColumn("timestamp_parsed4", to_timestamp("timestamp_str", "yyyy.MM.dd HH:mm:ss"))
```

Date to String Format

```

from pyspark.sql.functions import date_format
# 1. Format date as "yyyy-MM-dd"
# Input: 2025-01-25 (Date Type)
# Output: "2025-01-25" (String Type)
df = df.withColumn("formatted_date1", date_format("date_parsed", "yyyy-MM-dd"))

# 2. Format date as "dd-MMM-yyyy"
# Input: 2025-01-25 (Date Type)
# Output: "25-Jan-2025" (String Type)
df = df.withColumn("formatted_date2", date_format("date_parsed", "dd-MMM-yyyy"))

# 3. Format date as "MM/dd/yyyy"
# Input: 2025-01-25 (Date Type)
# Output: "01/25/2025" (String Type)
df = df.withColumn("formatted_date3", date_format("date_parsed", "MM/dd/yyyy"))

# 4. Format date as "dd/MM/yyyy"
# Input: 2025-01-25 (Date Type)
# Output: "25/01/2025" (String Type)
df = df.withColumn("formatted_date4", date_format("date_parsed", "dd/MM/yyyy"))

# 5. Format date as "MMMM dd, yyyy"
# Input: 2025-01-25 (Date Type)
# Output: "January 25, 2025" (String Type)
df = df.withColumn("formatted_date5", date_format("date_parsed", "MMMM dd, yyy"))

# 6. Format date as "EEE, dd MMM yyyy"
# Input: 2025-01-25 (Date Type)
# Output: "Sun, 25 Jan 2025" (String Type)
df = df.withColumn("formatted_date6", date_format("date_parsed", "EEE, dd MMM y"))

# 7. Format date as "yyyy/MM/dd"
# Input: 2025-01-25 (Date Type)
# Output: "2025/01/25" (String Type)
df = df.withColumn("formatted_date7", date_format("date_parsed", "yyyy/MM/dd"))

# 8. Format date as "yyyy.MM.dd"
# Input: 2025-01-25 (Date Type)
# Output: "2025.01.25" (String Type)
df = df.withColumn("formatted_date8", date_format("date_parsed", "yyyy.MM.dd"))

```

Timestamp to String Format

```
from pyspark.sql.functions import date_format
```

```
# 1. Format timestamp as "yyyy-MM-dd HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "2025-01-25 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp1", date_format("timestamp", "yyyy-MM-dd HH:mm:ss"))
```

```
# 2. Format timestamp as "dd-MMM-yyyy HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "25-Jan-2025 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp2", date_format("timestamp", "dd-MMM-yyyy HH:mm:ss"))
```

```
# 3. Format timestamp as "MM/dd/yyyy HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "01/25/2025 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp3", date_format("timestamp", "MM/dd/yyyy HH:mm:ss"))
```

```
# 4. Format timestamp as "dd/MM/yyyy HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "25/01/2025 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp4", date_format("timestamp", "dd/MM/yyyy HH:mm:ss"))
```

```
# 5. Format timestamp as "MMMM dd, yyyy HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "January 25, 2025 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp5", date_format("timestamp", "MMMM dd, yyyy HH:mm:ss"))
```

```
# 6. Format timestamp as "EEE, dd MMM yyyy HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "Sun, 25 Jan 2025 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp6", date_format("timestamp", "EEE, dd MMM yyyy HH:mm:ss"))
```

```
# 7. Format timestamp as "yyyy/MM/dd HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "2025/01/25 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp7", date_format("timestamp", "yyyy/MM/dd HH:mm:ss"))
```

```
# 8. Format timestamp as "yyyy.MM.dd HH:mm:ss"
```

```
# Input: "2025-01-25 10:15:00" (Timestamp Type)
```

```
# Output: "2025.01.25 10:15:00" (String Type)
```

```
df = df.withColumn("formatted_timestamp8", date_format("timestamp", "yyyy.MM.dd HH:mm:ss"))
```

```
# Show the resulting DataFrame
```

```
df.show(truncate=False)
```

Date Functions

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import (
    current_date, date_add, date_sub, datediff, add_months,
    trunc, date_format, year, month, dayofmonth, next_day, last_day
)

# 1. Current date
# "2025-01-25" -> Current system date (e.g., "2025-01-25")
df = df.withColumn("current_date", current_date())

# 2. Add 10 days to the date
# "2025-01-25" -> "2025-02-04"
df = df.withColumn("date_plus_10", date_add("date", 10))

# 3. Subtract 5 days from the date
# "2025-01-25" -> "2025-01-20"
df = df.withColumn("date_minus_5", date_sub("date", 5))

# 4. Difference in days from current date
# "2025-01-25" -> Number of days difference from today (e.g., "-5")
df = df.withColumn("days_diff", datediff(current_date(), "date"))

# 5. Add 2 months to the date
# "2025-01-25" -> "2025-03-25"
df = df.withColumn("add_months", add_months("date", 2))

# 6. Extract year
# "2025-01-25" -> "2025"
df = df.withColumn("year", year("date"))

# 7. Extract month
# "2025-01-25" -> "1"
df = df.withColumn("month", month("date"))

# 8. Extract day of the month
# "2025-01-25" -> "25"
df = df.withColumn("day", dayofmonth("date"))

# 9. Extract day of the week (1 = Sunday, 7 = Saturday)
# Input: "2025-01-25"
# Output: 7 (Saturday)
df = df.withColumn("day_of_week", dayofweek("date"))

# 10. Extract week of the year
# Input: "2025-01-25"
```

```
# Output: 4 (Week 4 of the year)
df = df.withColumn("week_of_year", weekofyear("date"))

# 11. Truncate to the first day of the month
# "2025-01-25" -> "2025-01-01"
df = df.withColumn("trunc_month", trunc("date", "MM"))

# 12. Next specified day of the week
# "2025-01-25" -> Next Monday (e.g., "2025-01-27")
df = df.withColumn("next_monday", next_day("date", "Monday"))

# 13. Last day of the month
# "2025-01-25" -> "2025-01-31"
df = df.withColumn("last_day_month", last_day("date"))
```

Time Functions

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import (
    current_timestamp, hour, minute, second, unix_timestamp, from_unixtime
)

# 1. Current timestamp
# "2025-01-25 10:15:00" -> Current system timestamp (e.g., "2025-01-25 10:15:00")
df = df.withColumn("current_timestamp", current_timestamp())

# 2. Extract hour
# "2025-01-25 10:15:00" -> "10"
df = df.withColumn("hour", hour("timestamp"))

# 3. Extract minute
# "2025-01-25 10:15:00" -> "15"
df = df.withColumn("minute", minute("timestamp"))

# 4. Extract second
# "2025-01-25 10:15:00" -> "00"
df = df.withColumn("second", second("timestamp"))

# 5. Convert date to Unix timestamp
# "2025-01-25 10:15:00" -> "1737763200"
df = df.withColumn("unix_timestamp", unix_timestamp("timestamp"))

# 6. Convert Unix timestamp to readable date
# "1737763200" -> "2025-01-25 10:15:00"
df = df.withColumn("from_unix", from_unixtime(unix_timestamp("timestamp")))
```

Math Functions

Simple Arithmetic

```
# 1. Add two columns  
# Input: col1 = 10, col2 = 5  
# Output: 15 (col1 + col2)  
df = df.withColumn("sum", col("col1") + col("col2"))  
  
# 2. Subtract two columns  
# Input: col1 = 10, col2 = 5  
# Output: 5 (col1 - col2)  
df = df.withColumn("difference", col("col1") - col("col2"))  
  
# 3. Multiply two columns  
# Input: col1 = 10, col2 = 5  
# Output: 50 (col1 * col2)  
df = df.withColumn("product", col("col1") * col("col2"))  
  
# 4. Divide two columns  
# Input: col1 = 10, col2 = 5  
# Output: 2.0 (col1 / col2)  
df = df.withColumn("quotient", col("col1") / col("col2"))  
  
# 5. Add a constant to a column  
# Input: col1 = 10  
# Output: 15 (col1 + 5)  
df = df.withColumn("sum_with_constant", col("col1") + 5)  
  
# 6. Subtract a constant from a column  
# Input: col1 = 10  
# Output: 5 (col1 - 5)  
df = df.withColumn("difference_with_constant", col("col1") - 5)
```

Complex Arithmetic


```
from pyspark.sql.functions import (
    abs, round, floor, ceil, exp, log, sqrt, pow
)

# 1. Absolute value
# Input: -2.71
# Output: 2.71
df = df.withColumn("abs_value", abs("value"))

# 2. Round the number to 2 decimal places
# Input: 3.14159
# Output: 3.14
df = df.withColumn("rounded_value", round("value", 2))

# 3. Floor (round down to the nearest integer)
# Input: 3.14
# Output: 3
df = df.withColumn("floor_value", floor("value"))

# 4. Ceil (round up to the nearest integer)
# Input: 3.14
# Output: 4
df = df.withColumn("ceil_value", ceil("value"))

# 5. Exponent (e raised to the power of the value)
# Input: 2.0
# Output: 7.389056
df = df.withColumn("exp_value", exp("value"))

# 6. Logarithm (log base e of the value)
# Input: 2.718
# Output: 0.999896
df = df.withColumn("log_value", log("value"))

# 7. Square root
# Input: 16
# Output: 4
df = df.withColumn("sqrt_value", sqrt("value"))

# 8. Power (raise the value to the power of 2)
# Input: 3
# Output: 9
df = df.withColumn("pow_value", pow("value", 2))
```

String Functions

Basic String Functions

1. Concatenate two strings

Input: "hello world" + " !!!"

Output: "hello world !!!"

```
df = df.withColumn("concatenated_2_cols", concat(col("col1"), col("col2")))
```

```
df = df.withColumn("concatenated_col_with_lit", concat(col("text"), lit(" !!!")))
```

2. Concatenate columns with a separator (Space)

Input: ("John", "Doe", "30")

Output: "John Doe 30"

```
df = df.withColumn("full_name", concat_ws(" ", col("first_name"), col("last_name"), col("age")))
```

3. Concatenate columns with a separator (Comma)

Input: ("John", "Doe", "30")

Output: "John, Doe, 30"

```
df = df.withColumn("full_name_comma", concat_ws(",", col("first_name"), col("last_name"), col("age")))
```

4. Concatenate with a custom string

Input: ("John", "Doe")

Output: "Name: John Doe"

```
df = df.withColumn("name", concat_ws("", lit("Name: "), col("first_name"), lit(" "), col("last_name")))
```

5. Check if string contains a substring

Input: "hello world" -> "world"

Output: True

```
df = df.withColumn("contains_world", col("text").contains("world"))
```

```
df = df.withColumn("contains_world2", contains(col("text"), lit("world")))
```

6. Check if string starts with a specific substring

Input: "hello world" -> "hello"

Output: True

```
df = df.withColumn("starts_with_hello", col("text").startswith("hello"))
```

```
df = df.withColumn("starts_with_hello2", startswith(col("text"), lit("hello")))
```

7. Check if string ends with a specific substring

Input: "hello world" -> "world"

Output: True

```
df = df.withColumn("ends_with_world", col("text").endswith("world"))
```

```
df = df.withColumn("ends_with_world2", endswith(col("text"), lit("world")))
```

8. Capitalize the first letter of each word

Input: "hello world"

```
# Output: "Hello World"
df = df.withColumn("initcap_text", initcap(col("text")))

# 9. Convert string to uppercase
# Input: "hello world"
# Output: "HELLO WORLD"
df = df.withColumn("upper_text", upper(col("text")))

# 10. Convert string to lowercase
# Input: "HELLO WORLD"
# Output: "hello world"
df = df.withColumn("lower_text", lower(col("text")))

# 11. Get the length of the string
# Input: "hello world"
# Output: 11
df = df.withColumn("length_of_text", length(col("text")))
```

Trim and Pad Functions

```
# 1. Trim: Remove both leading and trailing spaces from first_name
# Input: " John "
# Output: "John"
df = df.withColumn("trimmed_first_name", trim(col("first_name")))

# 2. Ltrim: Remove leading spaces from first_name
# Input: " John"
# Output: "John"
df = df.withColumn("ltrim_first_name", ltrim(col("first_name")))

# 3. Rtrim: Remove trailing spaces from last_name
# Input: "Doe "
# Output: "Doe"
df = df.withColumn("rtrim_last_name", rtrim(col("last_name")))

# 4. Lpad: Pad first_name with spaces on the left to make the length 10
# Input: "John"
# Output: "    John"
df = df.withColumn("lpad_first_name", lpad(col("first_name"), 10, " "))

# 5. Rpad: Pad last_name with spaces on the right to make the length 10
# Input: "Doe"
# Output: "Doe    "
df = df.withColumn("rpad_last_name", rpad(col("last_name"), 10, " "))

# 6. Lpad with a custom padding character: Pad first_name with "0" on the left to make the length 10
# Input: "John"
# Output: "00000John"
df = df.withColumn("lpad_first_name_zeros", lpad(col("first_name"), 10, "0"))

# 7. Rpad with a custom padding character: Pad last_name with "0" on the right to make the length 10
# Input: "Doe"
# Output: "Doe00000000"
df = df.withColumn("rpad_last_name_zeros", rpad(col("last_name"), 10, "0"))
```

Advanced String Functions

```
# 1. Substring: Extract substring from the full_name starting from position 1 (inclusive)
# Input: "John_Doe_30"
# Output: "John"
df = df.withColumn("substring_example", substring(col("full_name"), 1, 4))

# 2. Substring: Extract substring from the full_name starting from position 6 (inclusive)
# Input: "John_Doe_30"
# Output: "Doe"
df = df.withColumn("substring_name", substring(col("full_name"), 6, 3))

# 3. Substring: Extract last 2 characters of the full_name
# Input: "John_Doe_30"
# Output: "30"
df = df.withColumn("substring_age", substring(col("full_name"), -2, 2))

# 4. Split: Split the full_name into first and last names based on the "_" separator
# Input: "John_Doe_30"
# Output: ["John", "Doe", "30"]
df = df.withColumn("split_name", split(col("full_name"), "_"))

# 5. Split: Split the full_name into first and last names based on the "_" separator and get the first part (first name)
# Input: "John_Doe_30"
# Output: "John"
df = df.withColumn("first_name", split(col("full_name"), "_")[0])

# 6. Split: Split the full_name into first and last names and get the second part (last name)
# Input: "John_Doe_30"
# Output: "Doe"
df = df.withColumn("last_name", split(col("full_name"), "_")[1])

# 7. Split: Split the full_name and get the third part (age)
# Input: "John_Doe_30"
# Output: "30"
df = df.withColumn("age", split(col("full_name"), "_")[2])
```

Converting to Other Data Types

```
# 1. Convert string to integer
# Input: "12345"
# Output: 12345 (as an Integer type)
df = df.withColumn("int_parsed", col("int_str").cast("int"))

# 2. Convert string to float
# Input: "123.45"
# Output: 123.45 (as a Float type)
df = df.withColumn("float_parsed", col("int_str").cast("float"))

# 3. Convert string to double
# Input: "123.4567"
# Output: 123.4567 (as a Double type)
df = df.withColumn("double_parsed", col("int_str").cast("double"))

# 4. Convert string to long
# Input: "123456789012"
# Output: 123456789012 (as a Long type)
df = df.withColumn("long_parsed", col("int_str").cast("long"))

# 5. Convert integer to string
# Input: 12345
# Output: "12345" (as a String type)
df = df.withColumn("int_to_str", col("int_parsed").cast("string"))

# 6. Convert date to string
# Input: 2025-01-25 (Date type)
# Output: "2025-01-25" (String type)
df = df.withColumn("date_to_str", col("date_parsed").cast("string"))

# 7. Convert timestamp to string
# Input: 2025-01-25 10:15:00 (Timestamp type)
# Output: "2025-01-25 10:15:00" (String type)
df = df.withColumn("timestamp_to_str", col("timestamp_parsed").cast("string"))
```

Window Functions

Basic Window Functions

Use orderBy() when order matters:

- Ranking Functions (row_number, rank, dense_rank)
- Offset Functions (lead, lag)
- Cumulative Aggregations (sum, avg with rowsBetween)

Skip orderBy() when order is irrelevant:

- Partition-wise Aggregates (sum, avg, count)
- Row-Agnostic Aggregations (max, min)

```
from pyspark.sql.window import Window
from pyspark.sql.functions import col, row_number, rank, dense_rank, lag, lead

# Define window specification (partition by department, order by salary descending)
window_spec = Window.partitionBy("department").orderBy(col("salary").desc())

# Apply window functions

##row_number:** Assigns unique numbers to each row in a partition.
df = df.withColumn("row_number", row_number().over(window_spec))

##rank:** Similar to row_number but allows rank gaps.
df = df.withColumn("rank", rank().over(window_spec))

##dense_rank:** Like rank but without gaps.
df = df.withColumn("dense_rank", dense_rank().over(window_spec))

##lag:** Gets the previous row's value.
df = df.withColumn("previous_salary", lag("salary").over(window_spec))

##lead:** Gets the next row's value.
df = df.withColumn("next_salary", lead("salary").over(window_spec))

##sum:** Computes a running total.
df = df.withColumn("running_total", sum("salary").over(window_spec))

##avg:** Computes a moving average.
df = df.withColumn("moving_avg", avg("salary").over(window_spec))

# Show result
df.show()
```

With Rows Between

```
from pyspark.sql.window import Window
from pyspark.sql.functions import col, sum, avg, min, max, count

#1. Rolling sum over the last 2 rows and current row
window_spec1 = Window.partitionBy("department").orderBy("salary").rowsBetween
df = df.withColumn("rolling_sum_last_2", sum("salary").over(window_spec1))

#2. Moving average including previous, current, and next row
window_spec2 = Window.partitionBy("department").orderBy("salary").rowsBetween
df = df.withColumn("moving_avg", avg("salary").over(window_spec2))

#3. Rolling minimum for current and next 2 rows
window_spec3 = Window.partitionBy("department").orderBy("salary").rowsBetween
df = df.withColumn("rolling_min_next_2", min("salary").over(window_spec3))

#4. Maximum salary over all previous rows (running max)
window_spec4 = Window.partitionBy("department").orderBy("salary").rowsBetween
df = df.withColumn("running_max", max("salary").over(window_spec4))

#5. Count total rows within the window (entire partition)
window_spec5 = Window.partitionBy("department").orderBy("salary").rowsBetween
df = df.withColumn("total_rows", count("salary").over(window_spec5))

# Show result
df.show()
```

Array Functions

Creating and Manipulating Arrays


```
from pyspark.sql.functions import array, size, sort_array, array_contains, explode
```

```
# 1. Create an array column from multiple columns
```

```
# Input: col1 = "A", col2 = "B"
```

```
# Output: ["A", "B"]
```

```
df = df.withColumn("combined_array", array("col1", "col2"))
```

```
# 2. Get the size of an array
```

```
# Input: ["A", "B", "C"]
```

```
# Output: 3
```

```
df = df.withColumn("array_size", size("tags"))
```

```
# 3. Sort array elements in ascending order
```

```
# Input: [3, 1, 2]
```

```
# Output: [1, 2, 3]
```

```
df = df.withColumn("sorted_array", sort_array("numbers"))
```

```
# 4. Check if array contains a specific value
```

```
# Input: ["red", "blue"], check for "blue"
```

```
# Output: true
```

```
df = df.withColumn("has_blue", array_contains("colors", "blue"))
```

Array Elements

```
from pyspark.sql.functions import element_at, col
```

```
# 1. Get element at a specific position (1-based index)
```

```
# Input: ["x", "y", "z"], index = 2
```

```
# Output: "y"
```

```
df = df.withColumn("second_item", element_at("letters", 2))
```

```
# 2. Get element using 0-based index (like Python)
```

```
# Input: ["x", "y", "z"], index = 1
```

```
# Output: "y"
```

```
df = df.withColumn("second_item_alt", col("letters").getItem(1))
```

Modifying Array Elements

```
from pyspark.sql.functions import array_remove, array_distinct, array_union
```

```
# 1. Remove a specific element from an array
```

```
# Input: [1, 2, 2, 3], remove 2
```

```
# Output: [1, 3]
```

```
df = df.withColumn("no_twos", array_remove("numbers", 2))
```

```
# 2. Remove duplicate elements from array
```

```
# Input: [1, 2, 2, 3]
```

```
# Output: [1, 2, 3]
```

```
df = df.withColumn("unique_numbers", array_distinct("numbers"))
```

```
# 3. Merge two arrays and remove duplicates
```

```
# Input: [1, 2], [2, 3]
```

```
# Output: [1, 2, 3]
```

```
df = df.withColumn("merged_array", array_union("arr1", "arr2"))
```

Arrays to Rows

```
from pyspark.sql.functions import explode, posexplode
```

```
# 1. Convert array elements into multiple rows
```

```
# Input: ["apple", "banana"]
```

```
# Output: 2 rows: "apple", "banana"
```

```
df = df.withColumn("fruit", explode("fruits"))
```

```
# 2. Explode with position (index)
```

```
# Input: ["apple", "banana"]
```

```
# Output: (0, "apple"), (1, "banana")
```

```
df = df.select("id", posexplode("fruits").alias("pos", "fruit"))
```

Rows to Array

```
from pyspark.sql.functions import collect_list, collect_set
```

```
# Input: Multiple rows with same category, e.g. {"cat": "A", "val": 1}, {"cat": "A", "val": 2}  
# Output: {"cat": "A", "val_array": [1, 2]}
```

```
# 1. Group rows into array (with duplicates)
```

```
df_grouped = df.groupBy("cat").agg(collect_list("val").alias("val_array"))
```

```
# 2. Group rows into array (unique values only)
```

```
df_grouped = df.groupBy("cat").agg(collect_set("val").alias("unique_vals"))
```

Running SQL Queries

With Temp View

```
# Create a temporary SQL table from a DataFrame
df.createOrReplaceTempView("employees")

# Select all columns
df_sql = spark.sql("SELECT * FROM employees")

# Select specific columns
df_sql = spark.sql("SELECT name, salary FROM employees")

# Filter data
df_sql = spark.sql("SELECT * FROM employees WHERE salary > 50000")

# Aggregations
df_sql = spark.sql("SELECT department, AVG(salary) AS avg_salary FROM employees")

# Sorting
df_sql = spark.sql("SELECT * FROM employees ORDER BY salary DESC")

# Using LIMIT
df_sql = spark.sql("SELECT * FROM employees LIMIT 10")

# Using CASE WHEN
df_sql = spark.sql("""
    SELECT name, salary,
    CASE
        WHEN salary > 50000 THEN 'High'
        ELSE 'Low'
    END AS salary_category
    FROM employees
""")
```

Without Temp View

```
# Load any dataframe
df = spark.read.format('csv').option('header', 'true').load('/samples/customers.csv')
# Use Spark SQL with a variable and pass the dataframe
spark.sql("select first_name from {customers_df}", customers_df = df).show()

# Load any dataframe
df2 = spark.read.format('csv').option('header', 'true').load('/samples/orders.csv')
# Use Spark SQL with a variable and pass the dataframe
spark.sql("select order_id from {orders_df}", orders_df = df2).show()
```